

BT2101 Tutorial 3

Ensemble Learning

Assignment 2

- Please set “shuffle=True”, “random_state” serves “shuffle” step

random_state : int, RandomState instance or None, optional, default=None

If int, random_state is the seed used by the random number generator; If RandomState instance, random_state is the random number generator; If None, the random number generator is the RandomState instance used by np.random. Used when `shuffle == True`.

- http://scikit-learn.org/stable/modules/generated/sklearn.model_selection.KFold.html

Assignment 2

- Please set “shuffle=True”, “random_state” serves “shuffle” step

Step 1. Do K-fold Cross Validation

- Set `n_splits=5` : 5-fold cross validation
- Set `random_state` to `12345`
- Set `shuffle` to `True`

Hint: Use `KFold` function

```
# Create a 5-fold cross validation  
kf = KFold(n_splits=5, shuffle=True, random_state=12345)
```

Assignment 2

Comparison between L1 (Lasso) and L2 (Ridge) regularization

L1 regularization on least squares:

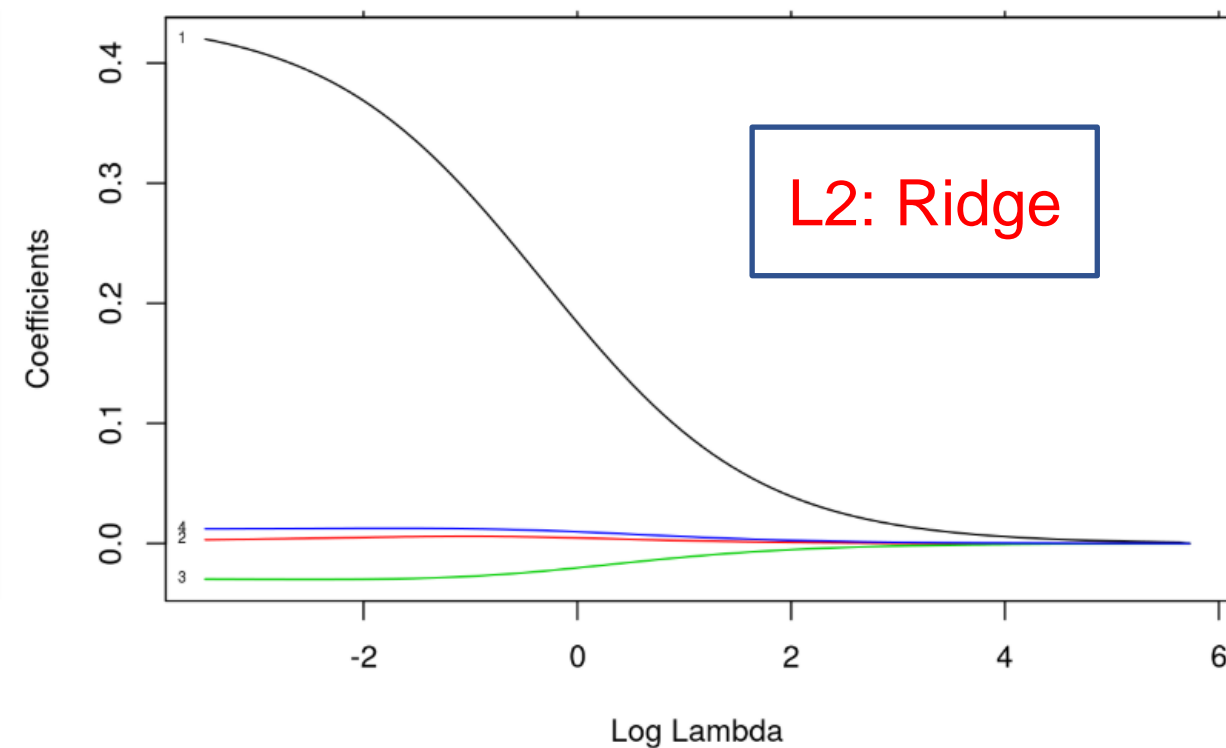
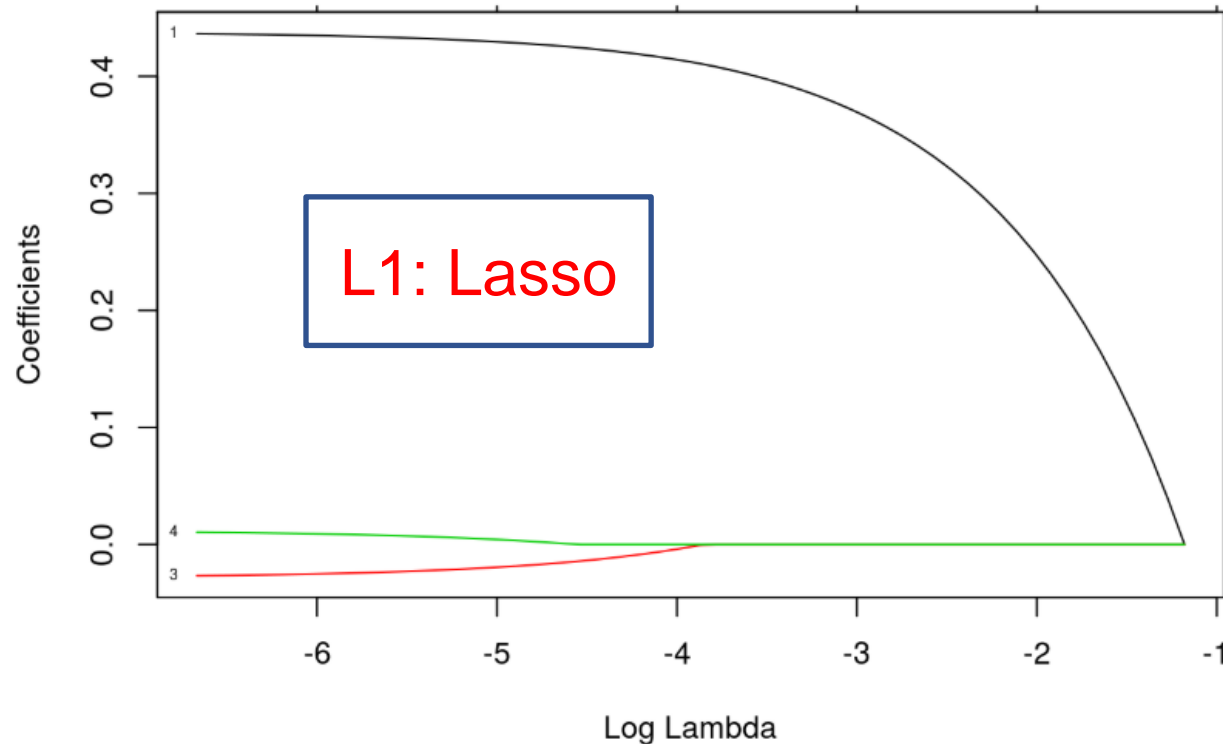
$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_j \left(t(\mathbf{x}_j) - \sum_i w_i h_i(\mathbf{x}_j) \right)^2 + \lambda \sum_{i=1}^k |w_i|$$

L2 regularization on least squares:

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_j \left(t(\mathbf{x}_j) - \sum_i w_i h_i(\mathbf{x}_j) \right)^2 + \lambda \sum_{i=1}^k w_i^2$$

Assignment 2

Comparison between L1 (Lasso) and L2 (Ridge) regularization



Assignment 2

Some online resources:

- <http://statweb.stanford.edu/~tibs/sta305files/Rudyregularization.pdf>
- <https://www.analyticsvidhya.com/blog/2016/01/complete-tutorial-ridge-lasso-regression-python/>
- http://scikit-learn.org/stable/auto_examples/linear_model/plot_logistic_l1_l2_sparsity.html
- <The Elements of Statistical Learning> Page 61-73
<https://web.stanford.edu/~hastie/Papers/ESLII.pdf>


Agenda

- Understand Ensemble Learning
- Discussion about Programming Assignment 3
 - BAGGING
 - Random Forest
 - AdaBoost
- Python Implementation

Ensemble Learning Method

What is Ensemble Method?

- Think about



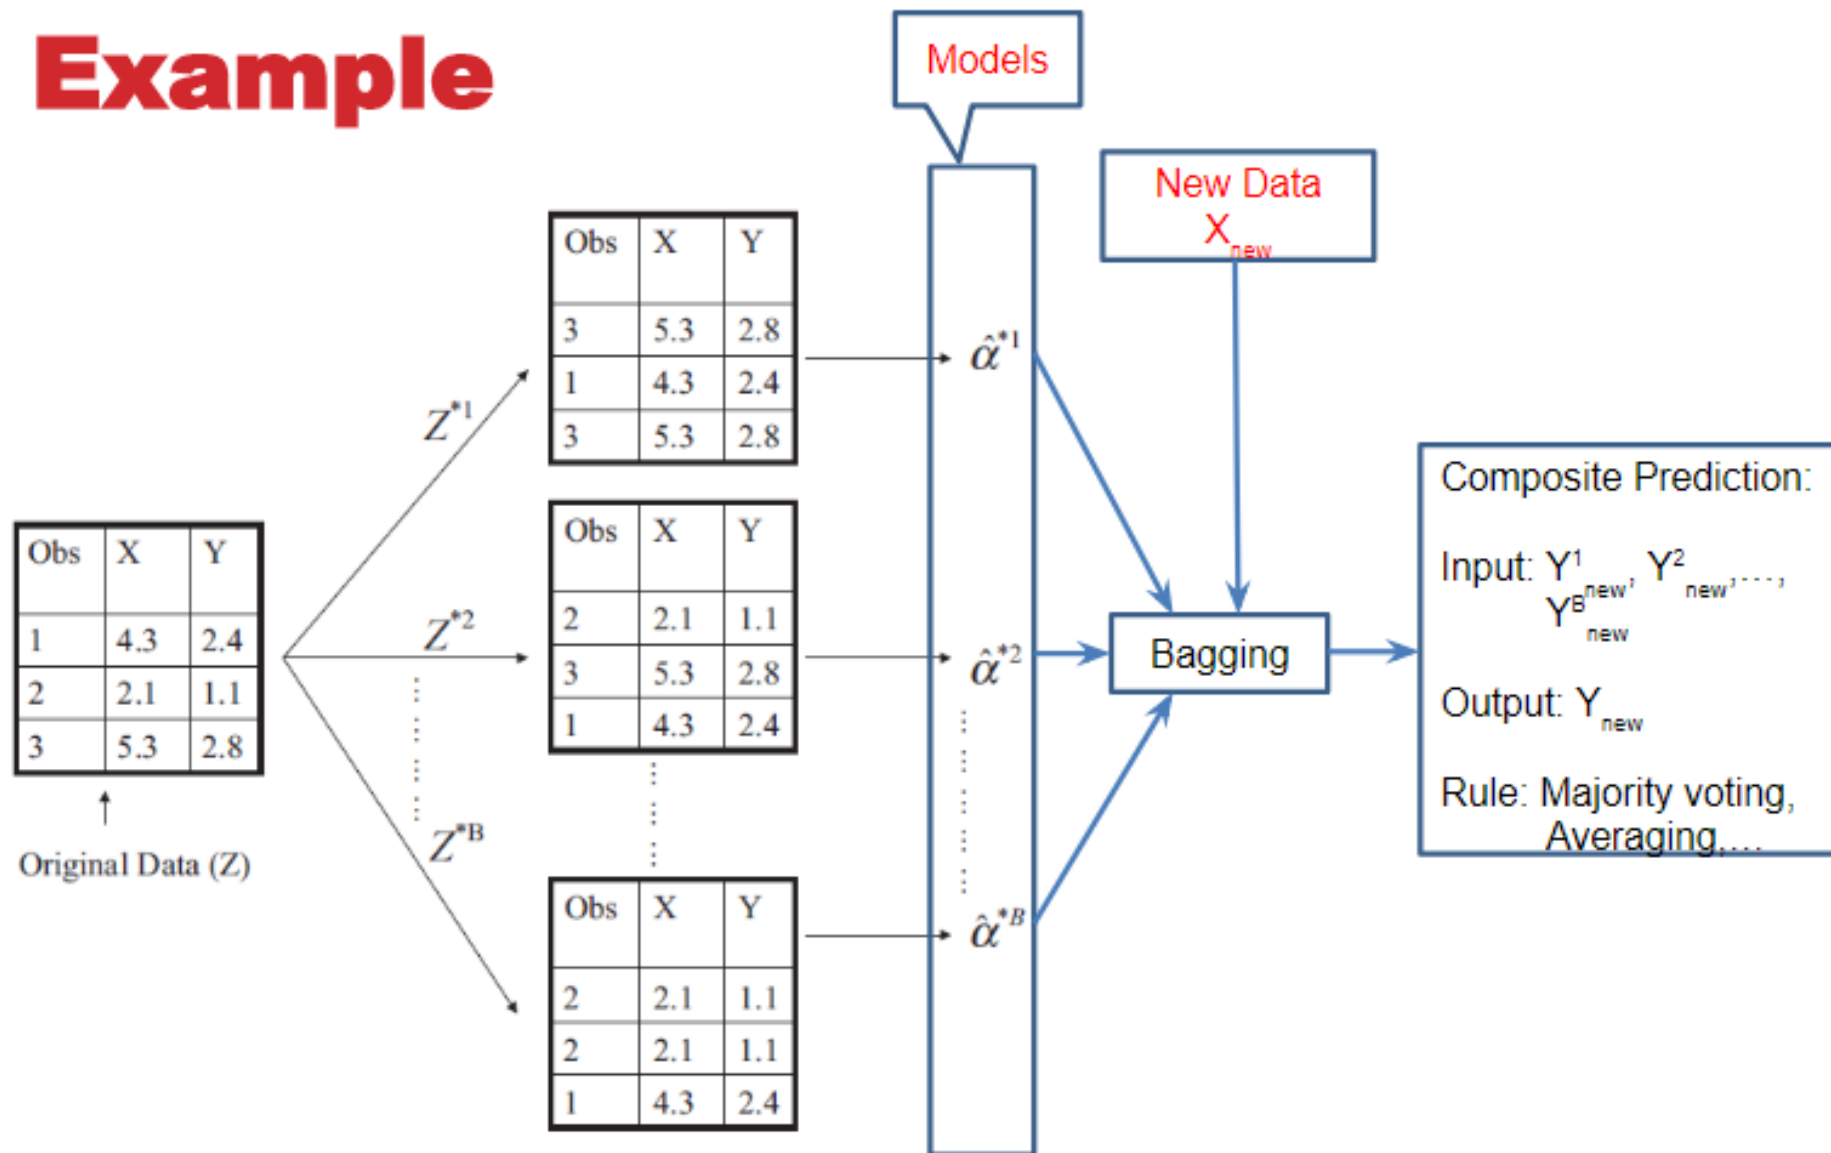
Learn from not just one classifier but a set of base classifiers, and combine their predictions for the classification of unseen instances using some form of voting (e.g., majority voting) (<Principles of Data Mining>)

Ensemble Method

- Bagging
- Random Forest
- Boosting
 - Adaboost
 - Gradient boosting
 - XGBoost
 - LightGBM
- Stacking

Bagging Example

- Bagging



Bagging

Example

Step 1.

Get your training dataset, suppose data size is N ($=3$ in this example)

Obs	X	Y
1	4.3	2.4
2	2.1	1.1
3	5.3	2.8

Original Data (Z)

Z^{*1}

Obs	X	Y
3	5.3	2.8
1	4.3	2.4
3	5.3	2.8

Z^{*2}

Obs	X	Y
2	2.1	1.1
3	5.3	2.8
1	4.3	2.4

Z^{*B}

Obs	X	Y
2	2.1	1.1
2	2.1	1.1
1	4.3	2.4

Models

$\hat{\alpha}^{*1}$

$\hat{\alpha}^{*2}$

$\hat{\alpha}^{*B}$

New Data
 X_{new}

Bagging

Composite Prediction:

Input: $Y_{\text{new}}^1, Y_{\text{new}}^2, \dots, Y_{\text{new}}^B$

Output: Y_{new}

Rule: Majority voting,
Averaging,...

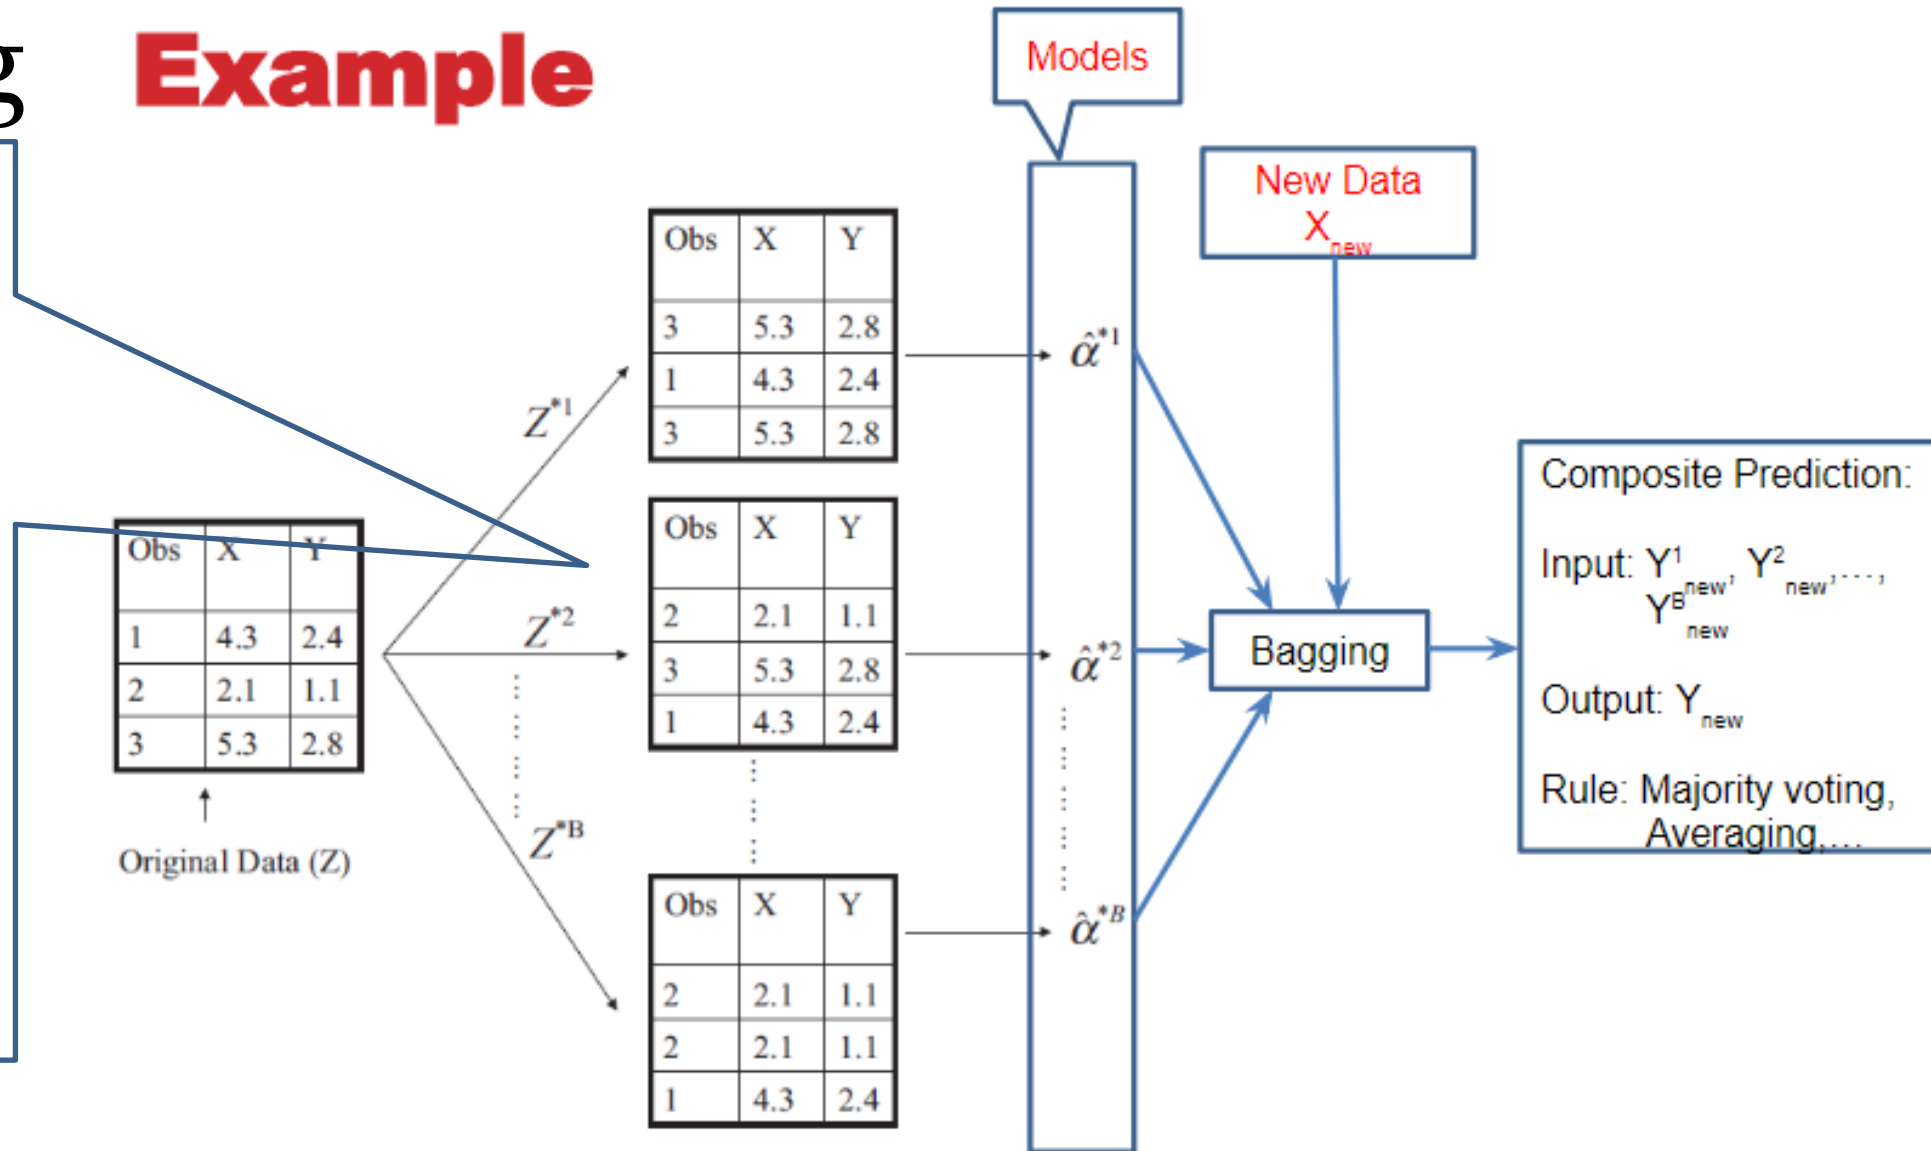
Bagging Example

Step 2.

Bootstrap sampling
(with replacement)
from the original
dataset B times

You get B bootstrap
samples

Each bootstrap
sample size is N
(usually same size
as original dataset)

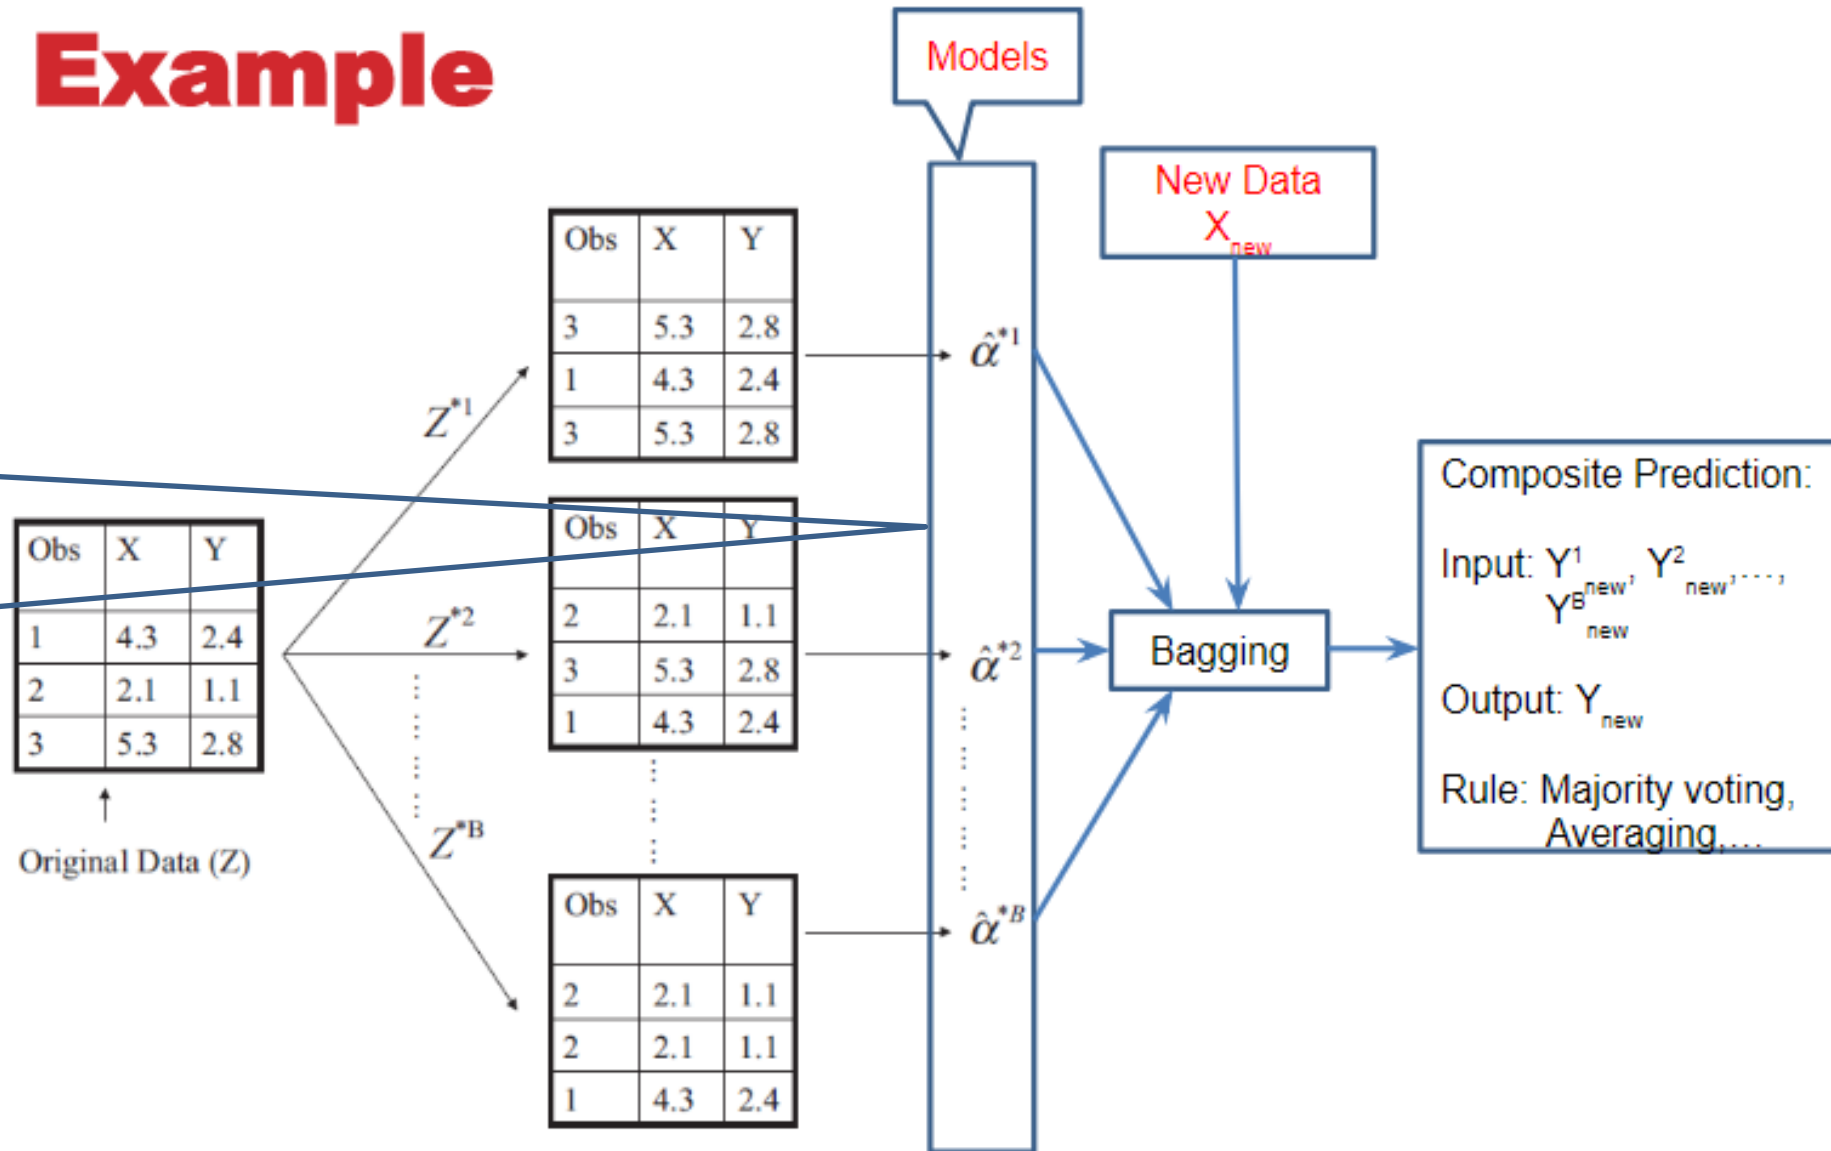


Bagging Example

Step 3.

You train B base models using these B bootstrap samples

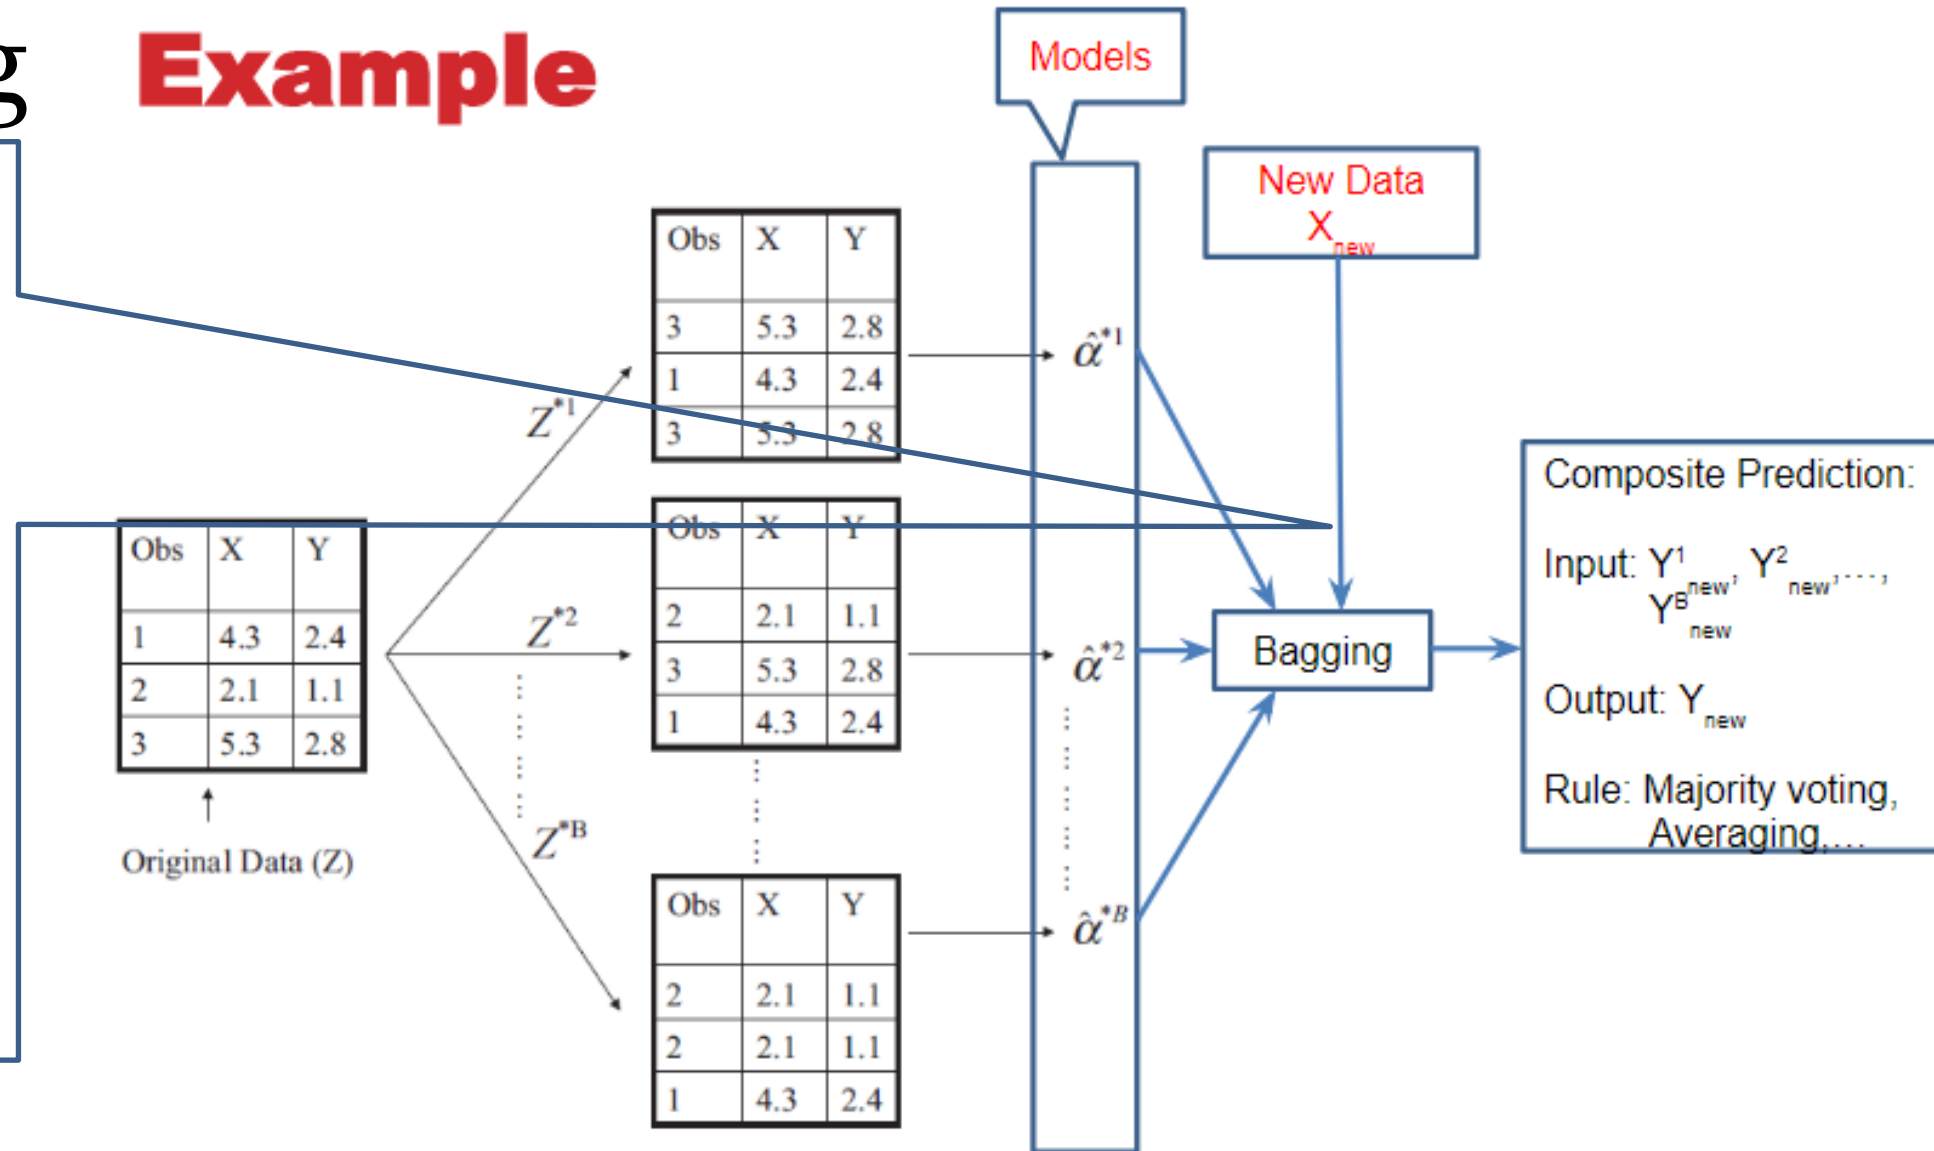
Now you have B base models



Bagging Example

Step 4.

Combine these base models into a “bag”



Bagging

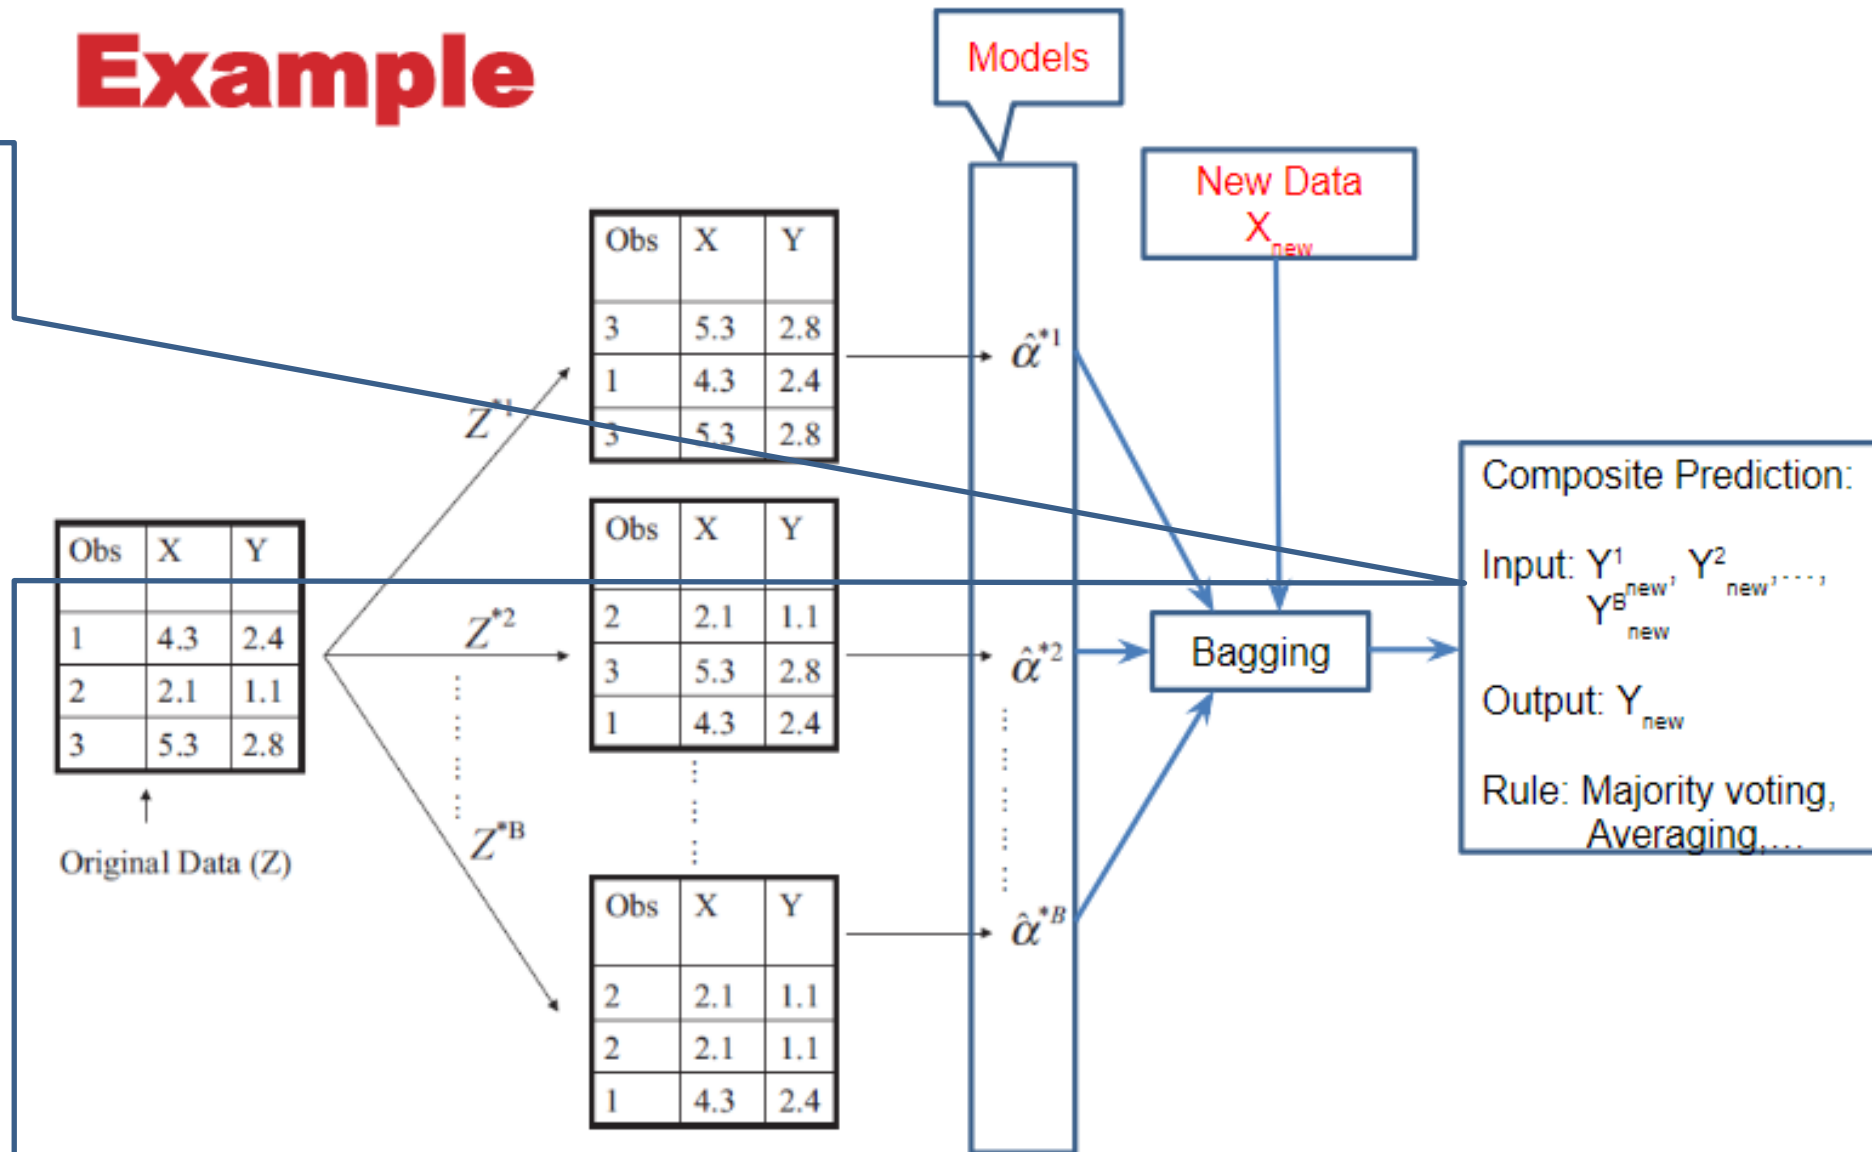
Example

Step 5.

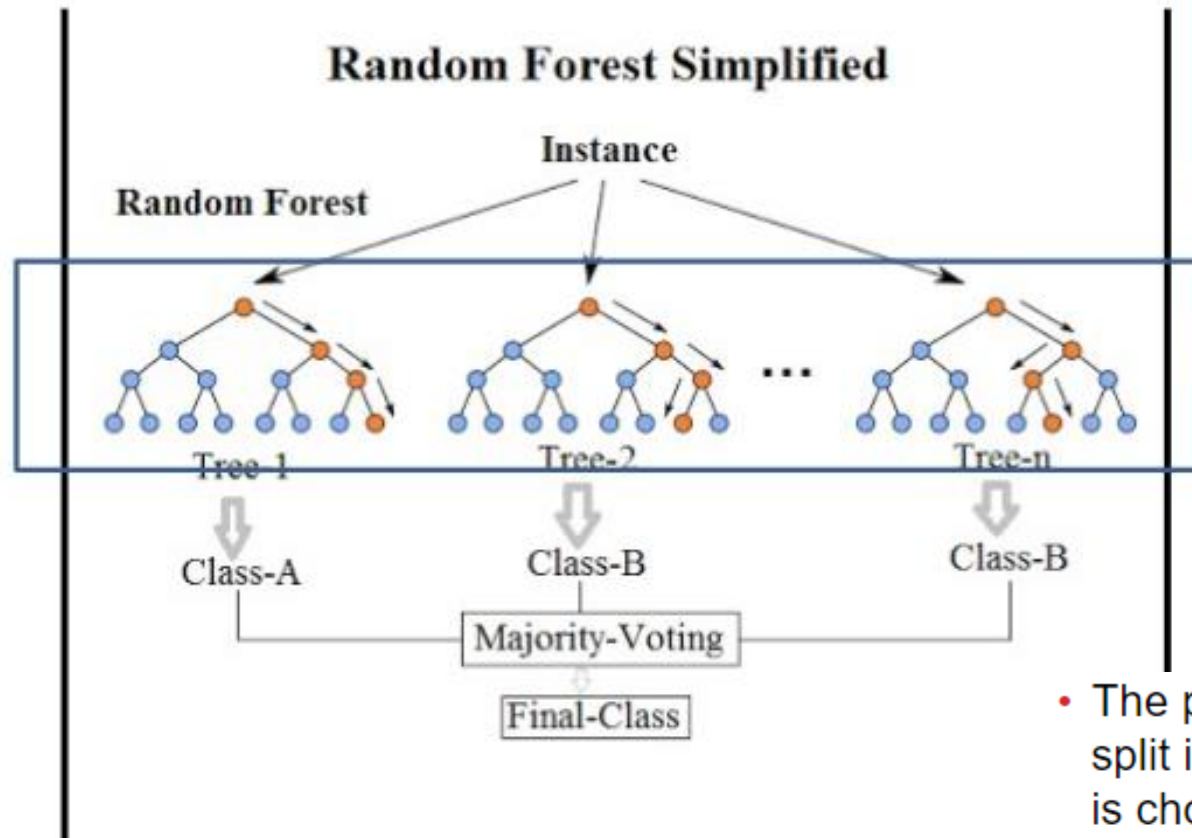
Suppose you want to predict output of new data sample X_{new}

Get predictions from B base models:
($Y^1_{\text{new}}, Y^2_{\text{new}}, \dots, Y^B_{\text{new}}$)

Composite prediction by averaging (for continuous output) or majority voting (for classification)



RANDOM FORESTS

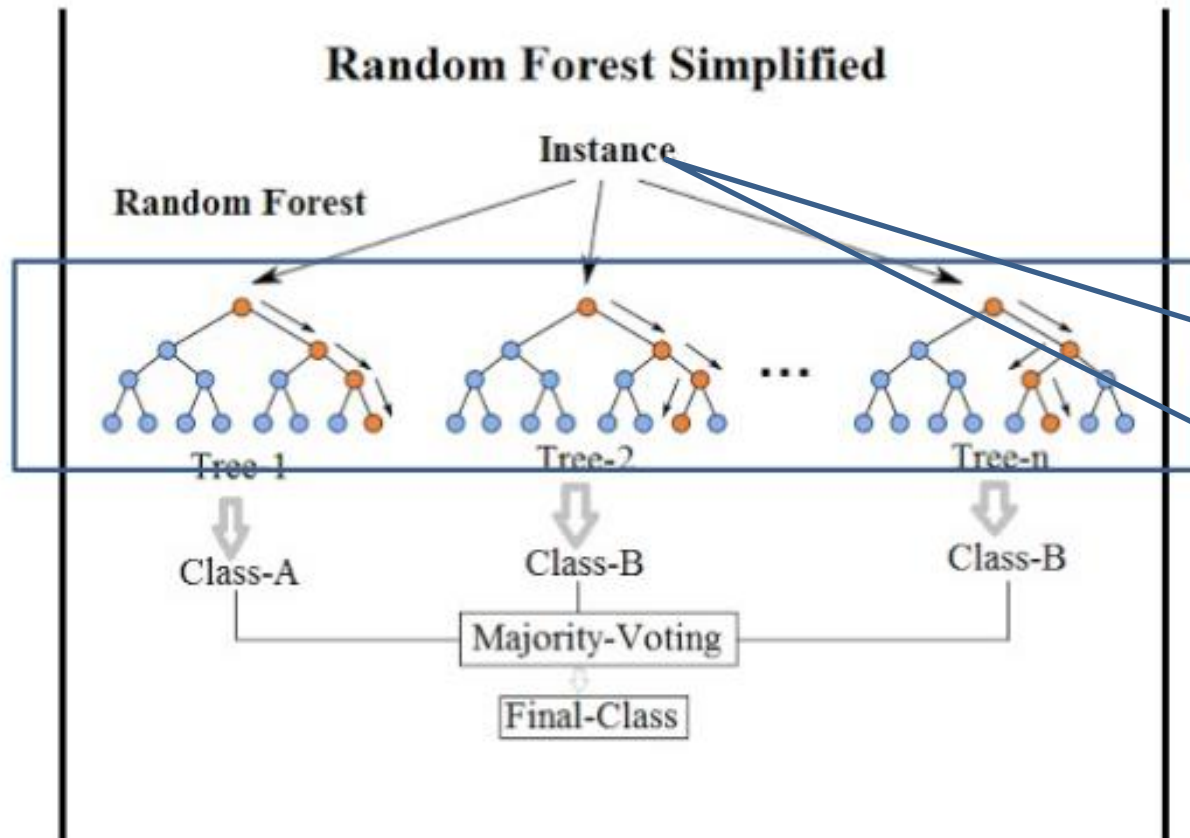


Models:

Randomly select
features to train
decision tree
models

- The process is similar to bagging except that, each time a split in a tree is considered, a random sample of m predictors is chosen as split candidate from the full set of p predictors
- In Practice, $m \approx \sqrt{p}$
- Bagging is a special case of random forest where $m = p$

RANDOM FORESTS



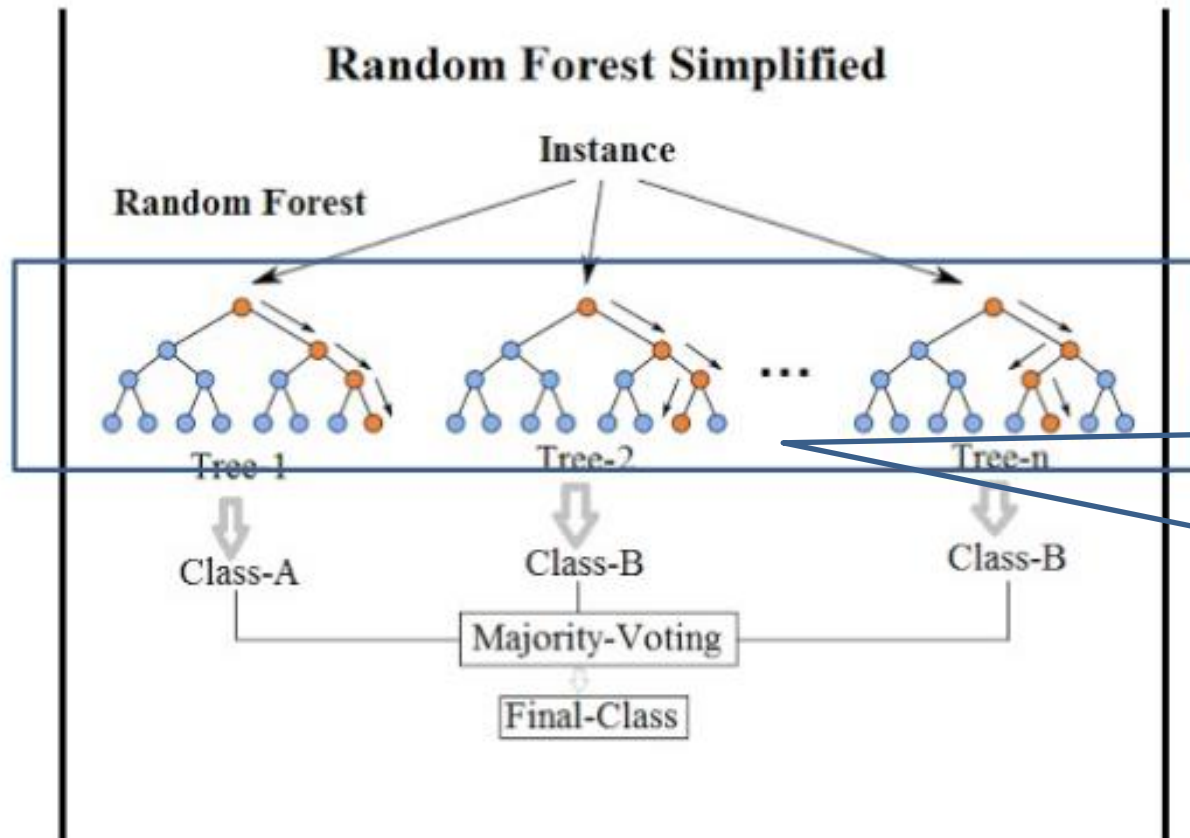
Models:

Randomly select
features to train
decision tree
models

Step 1.

Get your training
dataset, suppose
data size is N

RANDOM FORESTS



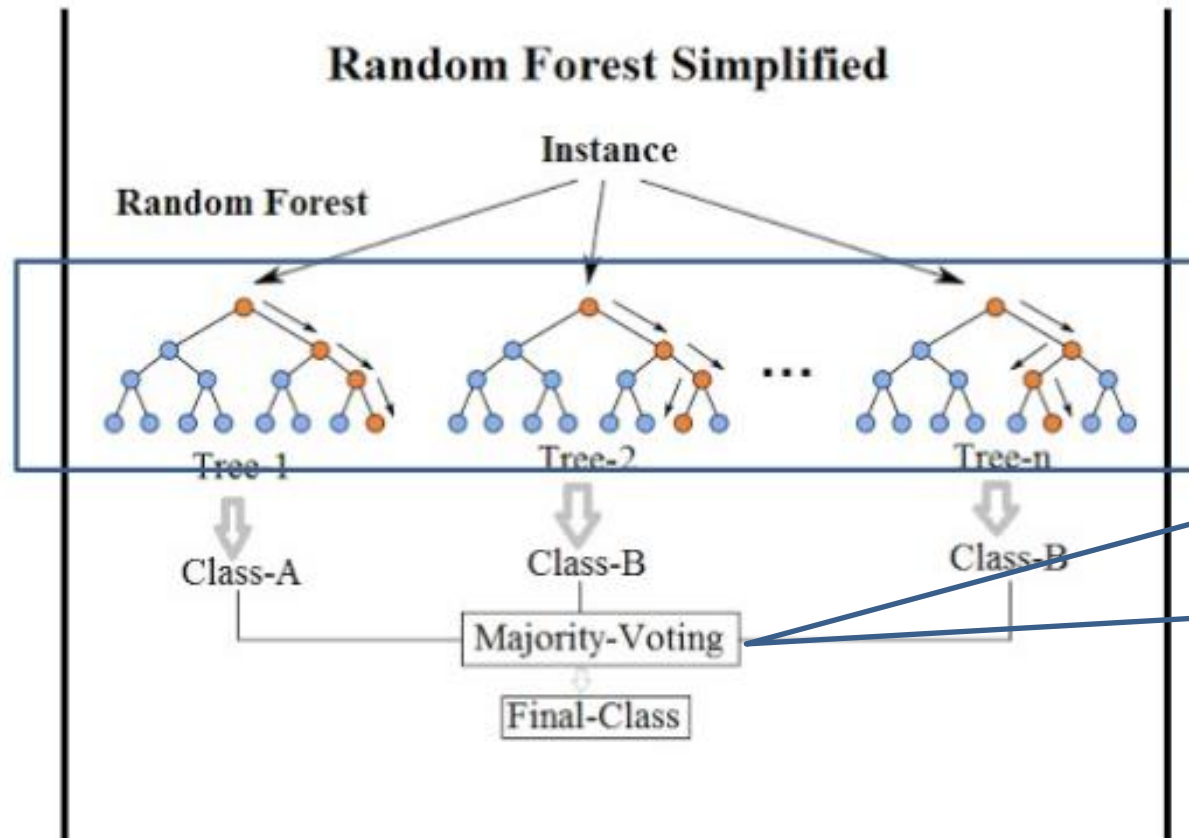
Models:

Randomly select
features to train
decision tree
models

Step 2.

Create n base models:
Each time, randomly select a
subset of original features to
get the base model

RANDOM FORESTS



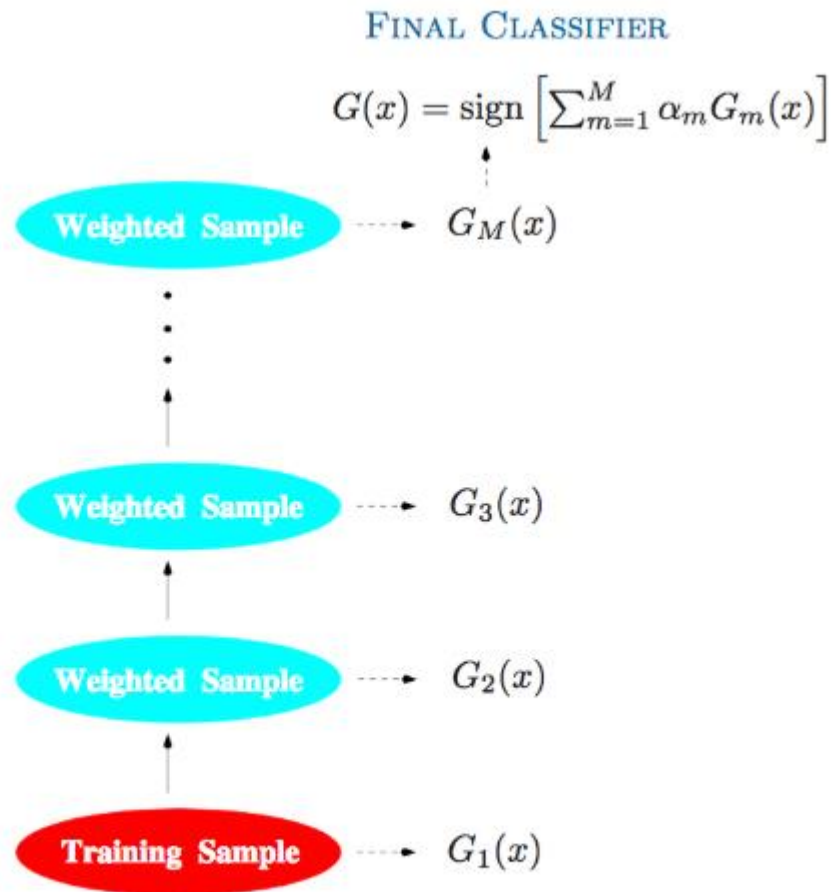
Models:

Randomly select
features to train
decision tree
models

Step 3.

Composite prediction:
Similar to bagging

Boosting Method



Difference?

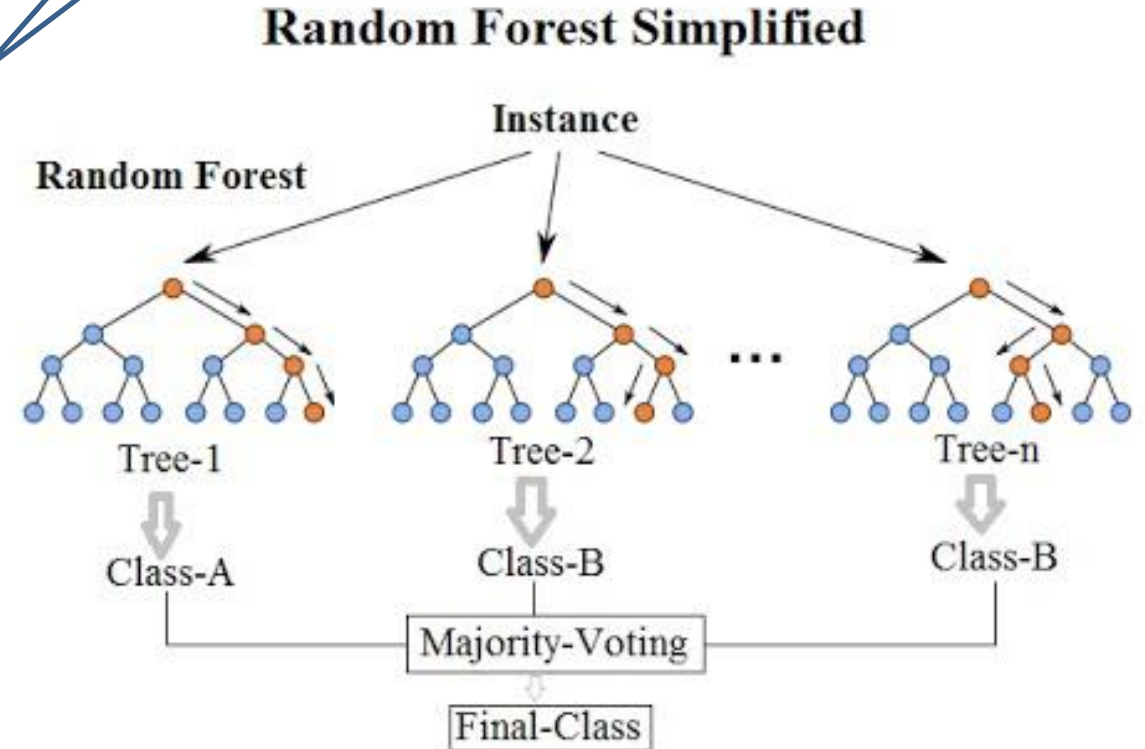


FIGURE 10.1. Schematic of AdaBoost. Classifiers are trained on weighted versions of the dataset, and then combined to produce a final prediction.

AdaBoost-Binary Classification

Algorithm 10.1 *AdaBoost.M1*.

1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.
2. For $m = 1$ to M :
 - (a) Fit a classifier $G_m(x)$ to the training data using weights w_i .
 - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
 - (c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.
 - (d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.
3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.

AdaBoost-Example for Step 2(b) to 2(d)

Algorithm 10.1 *AdaBoost.M1*.

1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.
 2. For $m = 1$ to M :
 - (a) Fit a classifier $G_m(x)$ to the training data using weights w_i .
 - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
 - (c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.
 - (d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.
 3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.
-

AdaBoost-Example for Step 2(b) to 2(d)

Row	x1	x2	x3	y	Weight ^(m-1)
1	0.77	0.35	0.31	1	0.2
2	0.87	0.78	0.59	1	0.2
3	0.34	0.32	0.23	-1	0.2
4	0.02	0.91	0.16	-1	0.2
5	0.61	0.55	0.04	1	0.2

Suppose we have this dataset: $N=5$

In step m , we train the dataset on y and get model $G_m(x)$

AdaBoost-Example for Step 2(b) to 2(d)

Row	x1	x2	x3	y	Weight ^(m-1)	G _m
1	0.77	0.35	0.31	1	0.2	1
2	0.87	0.78	0.59	1	0.2	1
3	0.34	0.32	0.23	-1	0.2	1
4	0.02	0.91	0.16	-1	0.2	1
5	0.61	0.55	0.04	1	0.2	1

Predicting output
using model
 $G_m(x)$

AdaBoost-Example for Step 2(b) to 2(d)

Row	x1	x2	x3	y	Weight ^(m-1)	G _m	I(y≠y ^(m))
1	0.77	0.35	0.31	1	0.2	1	0
2	0.87	0.78	0.59	1	0.2	1	0
3	0.34	0.32	0.23	-1	0.2	1	1
4	0.02	0.91	0.16	-1	0.2	1	1
5	0.61	0.55	0.04	1	0.2	1	0

(b) Compute

$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$

Compute $\text{err}_m=0.4$

AdaBoost-Example for Step 2(b) to 2(d)

Row	x1	x2	x3	y	Weight ^(m-1)	G _m	I(y≠y ^(m))	Weight ^(m)
1	0.77	0.35	0.31	1	0.2	1	0	0.2
2	0.87	0.78	0.59	1	0.2	1	0	0.2
3	0.34	0.32	0.23	-1	0.2	1	1	0.3
4	0.02	0.91	0.16	-1	0.2	1	1	0.3
5	0.61	0.55	0.04	1	0.2	1	0	0.2

(c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.

(d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.

Update weights;
If you do not
renormalize and sum
up to 1

AdaBoost-Example for Step 2(b) to 2(d)

(c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.

(d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.

Usually,


If $0 < \text{err}_m < 0.5$ and misclassified on i :

$$w_i \leftarrow w_i \times \frac{1 - \text{err}_m}{\text{err}_m}: w_i \text{ is increasing}$$

Unusually,

If $0.5 < \text{err}_m < 1$ and misclassified on i :

$$w_i \leftarrow w_i \times \frac{1 - \text{err}_m}{\text{err}_m}: w_i \text{ is decreasing}$$



Intuition: If $\text{err}_m > 0.5$, model error is even higher than “random guess” (50% accuracy/error), then do the opposite to whatever this model tells you.

AdaBoost-Example for Step 2(b) to 2(d)

Quiz: Do we need a model with high error rate (e.g., Error=90%)?

Suppose 3 Financial Analysts:

Analyst 1. Accuracy=95%, Error=5%

Analyst 2. Accuracy=51%, Error=49%

Analyst 3. Accuracy=5%, Error=95%

Do you think Analyst 3 is a good analyst or not?



AdaBoost-Example for Step 2(b) to 2(d)

Row	x1	x2	x3	y	Weight ^(m)
1	0.77	0.35	0.31	1	0.2
2	0.87	0.78	0.59	1	0.2
3	0.34	0.32	0.23	-1	0.3
4	0.02	0.91	0.16	-1	0.3
5	0.61	0.55	0.04	1	0.2

Goto step $m+1$;

In step $m+1$, we train and get new model $G_{m+1}(x)$, and new predictions G_{m+1} , and update weights;

Repeat the procedure until getting M models

AdaBoost-Example for Step 3

Algorithm 10.1 *AdaBoost.M1*.

1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.
2. For $m = 1$ to M :
 - (a) Fit a classifier $G_m(x)$ to the training data using weights w_i .
 - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
 - (c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.
 - (d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.

3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.

AdaBoost-Example for Step 3

Suppose you want to predict the output of a new data sample x_{new}

Model m	α_m	$G_m(x_{\text{new}})$: Prediction on this new data sample	$\alpha_m * G_m(x_{\text{new}})$
1	2	+1	+2
2	1	-1	-1
3	-2	+1	-2

3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.

The signum function of a real number x is defined as follows:

$$\text{sgn}(x) := \begin{cases} -1 & \text{if } x < 0, \\ 0 & \text{if } x = 0, \\ 1 & \text{if } x > 0. \end{cases}$$

Output of x_{new} : $\text{sign}(-1) = -1$

Intuition: If $\text{err}_m > 0.5$, model error is even higher than “random guess” (50% accuracy/error), then do the opposite to whatever this model tells you.

AdaBoost-Example for Step 2(a)

Algorithm 10.1 *AdaBoost.M1*.

1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.

2. For $m = 1$ to M :

(a) Fit a classifier $G_m(x)$ to the training data using weights w_i .

(b) Compute

$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$

(c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.

(d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.

3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.

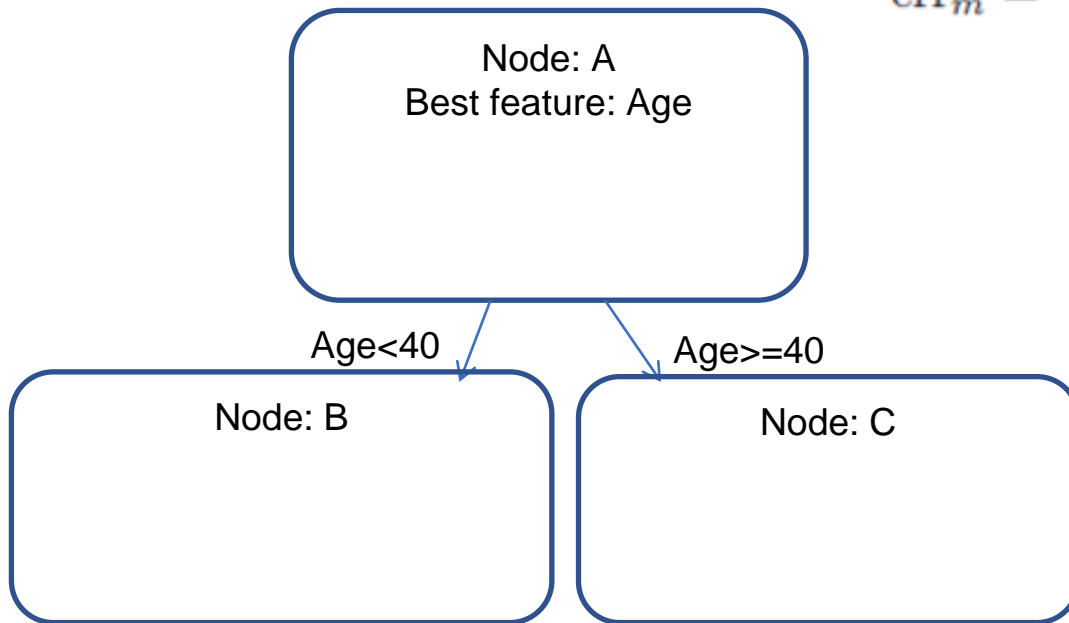
Select and extract a classifier G_m from the pool of classifiers, which minimizes err_m

$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$


AdaBoost-Example for Step 2(a)

Select and extract a classifier G_m from the pool of classifiers, which minimizes err_m

$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$



Model 1: err=40%



Model 2: err=10%

Which model do you choose in this step m?

AdaBoost-Binary Classification

Algorithm 10.1 *AdaBoost.M1*.

1. Initialize the observation weights $w_i = 1/N$, $i = 1, 2, \dots, N$.
 2. For $m = 1$ to M :
 - (a) Fit a classifier $G_m(x)$ to the training data using weights w_i .
 - (b) Compute
$$\text{err}_m = \frac{\sum_{i=1}^N w_i I(y_i \neq G_m(x_i))}{\sum_{i=1}^N w_i}.$$
 - (c) Compute $\alpha_m = \log((1 - \text{err}_m)/\text{err}_m)$.
 - (d) Set $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(x_i))]$, $i = 1, 2, \dots, N$.
 3. Output $G(x) = \text{sign} \left[\sum_{m=1}^M \alpha_m G_m(x) \right]$.
-

Limitation:

Adaboost may still have overfit issue.

Ensemble Performance

Performance of Binary Classifier

- Confusion Matrix


Correct classification	Classified as	
	+	−
+	true positives	false negatives
−	false positives	true negatives

True Positive Rate or Hit Rate or Recall or Sensitivity or TP Rate	TP/P	The proportion of positive instances that are correctly classified as positive
False Positive Rate or False Alarm Rate or FP Rate	FP/N	The proportion of negative instances that are erroneously classified as positive
False Negative Rate or FN Rate	FN/P	The proportion of positive instances that are erroneously classified as negative = $1 - \text{True Positive Rate}$

True Negative Rate or Specificity or TN Rate	TN/N	The proportion of negative instances that are correctly classified as negative
Precision or Positive Predictive Value	$TP / (TP + FP)$	Proportion of instances classified as positive that are really positive
F1 Score	$(2 \times \text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$	A measure that combines Precision and Recall
Accuracy or Predictive Accuracy	$(TP + TN) / (P + N)$	The proportion of instances that are correctly classified
Error Rate	$(FP + FN) / (P + N)$	The proportion of instances that are incorrectly classified

Cross Validation

- K-Fold Cross Validation (e.g., K=10)



- Model Evaluation
- Model Comparison
- Model Tuning

All observations are used for both training and validation, and each observation is used for validation exactly once.

Python Implementation

Implementation in Python

BT2101 Introduction to Ensemble Learning

1 Goal

In this notebook, we will explore **Ensemble Learning** including:

- Bagging
- Random Forest
- AdaBoost

For the **Decision Tree** method, you will:

- Use open-source package to do ensemble learning
- Compare performances of different methods

```
# -*- coding:utf-8 -*-  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from math import sqrt  
from __future__ import division  
%matplotlib inline
```

2 Ensemble Learning

Programming Assignment 3

Using the BT2101 Tutorial 3 Programming code ([Ensemble Learning.ipynb](#)), please answer the questions in the jupyter notebook

Answer all in the jupyter notebook.

Instructions

Submit Python Notebook to IVLE folder, and Naming the file:
AXXXX_T3_program.ipynb

Include your answers in the jupyter notebook

- You need to show outputs, instead of just showing functions.

Submit a **FINAL** program by **Tue Sep-18** (by lunchtime)

- Based on **Ensemble Learning.ipynb**

Thank you!

Appendix: AdaBoost-MultiClass Classification

Algorithm 1 *AdaBoost (Freund & Schapire 1997)*

1. Initialize the observation weights $w_i = 1/n$, $i = 1, 2, \dots, n$.

2. For $m = 1$ to M :

(a) Fit a classifier $T^{(m)}(\mathbf{x})$ to the training data using weights w_i .

(b) Compute

$$err^{(m)} = \sum_{i=1}^n w_i \mathbb{I}(c_i \neq T^{(m)}(\mathbf{x}_i)) / \sum_{i=1}^n w_i.$$

(c) Compute

$$\alpha^{(m)} = \log \frac{1 - err^{(m)}}{err^{(m)}}.$$

(d) Set

$$w_i \leftarrow w_i \cdot \exp \left(\alpha^{(m)} \cdot \mathbb{I}(c_i \neq T^{(m)}(\mathbf{x}_i)) \right), \quad i = 1, 2, \dots, n.$$

(e) Re-normalize w_i .

3. Output

$$C(\mathbf{x}) = \arg \max_k \sum_{m=1}^M \alpha^{(m)} \cdot \mathbb{I}(T^{(m)}(\mathbf{x}) = k).$$